

Lead Deduplication: Approaches Comparison and Analysis

Introduction: Deduplication Rules and Context

When working with a dataset of leads, duplicates need to be reconciled based on **certain rules**:

1. Duplicate Identification:

- Leads with the same `_id` are duplicates.
- Leads with the same email (case-insensitive) are also duplicates.
- Other fields don't determine duplication.

2. Resolution Strategy:

- The lead with the **latest** `entryDate` is preferred.
- If two leads share the **same entryDate**, the **last occurrence** in the list is preferred.

Approach 1: Simple Map-Based Deduplication

Approach 1 uses two hash maps to detect duplicates during a single pass:

- `idMap` keyed by `_id`
- `emailToldMap` keyed by normalized email

For each lead:

- If the `_id` or email is missing, it is classified as a bad lead.
- Check if the `_id` or email exists in the maps.
- If found, the existing and new lead are compared using a preference strategy (e.g., latest `entryDate` wins).
- The preferred lead replaces the existing one, maps updated accordingly.
- Field-level changes are logged.
- Otherwise, the lead is added as new.

This approach handles duplicates based on the presence of `_id` or email but **does not resolve transitive duplicates** where two leads might be linked through intermediate records (e.g., Lead A shares email with Lead B, Lead B shares `_id` with Lead C).

Algorithm:

1. Initialize empty `idMap`, `emailToldMap`, `badLeads` list.
2. For each lead:
 - Validate presence of `_id` and email; if missing, add to `badLeads`.
 - Normalize email.
 - Check for existing lead by `_id` or email.
 - If found, select preferred lead (using strategy), update maps.
 - Log changes if any differences detected.
 - If no existing lead found, add to maps.
3. Return deduplicated leads and bad leads.

Approach 2: Transitive Deduplication Using Union-Find (Disjoint Set)

How Transitive Deduplication Works:

For Example,

```
Lead a = {"_id": "1", "email": "x@example.com", "firstName": "A", "lastName": "X", "address": "Addr1", "entryDate": "2024-01-01T00:00:00Z"}
Lead b = {"_id": "2", "email": "x@example.com", "firstName": "B", "lastName": "Y", "address": "Addr2", "entryDate": "2024-01-02T00:00:00Z"}
Lead c = {"_id": "1", "email": "z@example.com", "firstName": "C", "lastName": "Z", "address": "Addr3", "entryDate": "2024-01-03T00:00:00Z"}
```

Let's Understand:

- a and b share email, so they are grouped.
- a and c share _id, so c joins the same group.
- Hence, a, b, c becomes 1 transitive group. (All the 3 Leads are Duplicates)
- Since Lead c, has the latest entryDate, it is selected as the preferred lead.

Approach 2 uses a **Union-Find (Disjoint Set)** data structure to detect clusters of duplicates connected via shared _id or email. It groups leads into sets where any two leads are linked transitively by common _id or email.

Steps:

- Separate out bad leads (missing _id or email).
- Assign each valid lead an index.
- For each lead:
 - Union leads sharing the same _id.
 - Union leads sharing the same email.
- After unions, leads belonging to the same root form a group.
- For each group:
 - Select the preferred lead (e.g., latest entryDate).
 - Log changes between preferred and other leads in the group.
 - Add preferred lead to final output maps.

This approach resolves **transitive duplicates** comprehensively.

Algorithm:

1. Filter out bad leads.
2. Initialize Union-Find data structure for all valid leads.
3. Map _id and email to lead indices, union leads with same _id or email.
4. Identify groups by finding roots in Union-Find.
5. For each group:
 - Determine preferred lead.
 - Log changes.
 - Add preferred lead to output maps.
6. Return deduplicated leads and bad leads.

Pros and Cons

Aspect	Approach 1: Simple Map-Based Deduplication	Approach 2: Union-Find (Transitive) Deduplication
Pros	Simple and straightforward to implement.	Detects all duplicates including transitive links.
	Efficient for datasets with direct _id or email duplicates.	More complete and accurate deduplication.
	Easy to maintain with minimal data structures.	Handles complex real-world data with overlapping identifiers.
Cons	cannot detect or merge transitive duplicates.	More complex to implement.
	May leave duplicates unresolved if identifiers are indirectly shared.	Slightly higher computational overhead due to union-find operations.
	Less robust for complex real-world data.	Requires extra memory to track union-find structures.

Algorithmic Complexity

Aspect	Approach 1	Approach 2 (Union-Find)
Time Complexity	$O(n)$ for scanning + $O(1)$ lookups per lead = $O(n)$	Union-Find with path compression = $O(n * \alpha(n))$ (very close to linear)
Space Complexity	$O(n)$ for storing maps and leads	$O(n)$ for storing parent maps and lead clusters

Where n is the number of leads, and $\alpha(n)$ is the inverse Ackermann function, practically ≤ 5 .

Conclusion: Why Approach 2 Was Selected

Approach 2 was selected for this project because:

- It **handles transitive duplicates** which are common in real-world datasets where duplicates share identifiers indirectly.
- Provides a **more comprehensive and reliable** deduplication.
- Ensures **no duplicates are left unresolved**, preventing data quality issues downstream.
- Though more complex, its additional robustness outweighs the minor performance costs.

Understanding Algorithms with Sample Input

Below is a sample input dataset illustrating lead records with overlapping _id and emails:

```
{
  "leads": [
    { "_id": "jkj238238jdsnfsj23", "email": "foo@bar.com", "firstName": "John", "lastName": "Smith", "address": "123 Street St", "entryDate": "2014-05-07T17:30:20+00:00" },
    { "_id": "edu45238jdsnfsj23", "email": "mae@bar.com", "firstName": "Ted", "lastName": "Masters", "address": "44 North Hampton St", "entryDate": "2014-05-07T17:31:20+00:00" },
    { "_id": "wabaj238238jdsnfsj23", "email": "bog@bar.com", "firstName": "Fran", "lastName": "Jones", "address": "8803 Dark St", "entryDate": "2014-05-07T17:31:20+00:00" },
    { "_id": "jkj238238jdsnfsj23", "email": "coo@bar.com", "firstName": "Ted", "lastName": "Jones", "address": "456 Neat St", "entryDate": "2014-05-07T17:32:20+00:00" },
    { "_id": "se1045238jdsnfsj23", "email": "foo@bar.com", "firstName": "John", "lastName": "Smith", "address": "123 Street St", "entryDate": "2014-05-07T17:32:20+00:00" },
    { "_id": "qest38238jdsnfsj23", "email": "foo@bar.com", "firstName": "John", "lastName": "Smith", "address": "123 Street St", "entryDate": "2014-05-07T17:32:20+00:00" },
    { "_id": "vug789238jdsnfsj23", "email": "fool@bar.com", "firstName": "Blake", "lastName": "Douglas", "address": "123 Reach St", "entryDate": "2014-05-07T17:33:20+00:00" },
    { "_id": "wuj08238jdsnfsj23", "email": "foo@bar.com", "firstName": "Micah", "lastName": "Valmer", "address": "123 Street St", "entryDate": "2014-05-07T17:33:20+00:00" },
    { "_id": "belr28238jdsnfsj23", "email": "mae@bar.com", "firstName": "Tallulah", "lastName": "Smith", "address": "123 Water St", "entryDate": "2014-05-07T17:33:20+00:00" },
    { "_id": "jkj238238jdsnfsj23", "email": "foo@bar.com", "firstName": "John", "lastName": "Smith", "address": "123 Street St", "entryDate": "2014-05-07T17:33:20+00:00" }
  ]
}
```

How Approach 1 Processes This Input

Step-by-Step:

1. Insert lead 1 → new _id, added.
2. Insert lead 2 → new _id, **but email matches** existing: maps email to existing _id=1, dedupes with lead 1.
3. Insert lead 1 again → _id already exists, dedupes.
4. Insert lead 3 → new _id, added.
5. Insert lead 4 → new _id, added.

Issues:

- Deduplication is local/pairwise.
- Email-based deduplication works only once (emailToIdMap can only point to one _id).
- Transitive chain 1 ↔ 2 ↔ 3 is **not fully resolved**.
- Retains 1, 3, 4 as separate leads.

How Approach 2 Processes This Input

Step-by-Step (Union-Find):

1. Insert 1 → own set
2. Insert 2 → union with 1 (via email)
3. Insert 1 again → already in group
4. Insert 3 → own group
5. Insert 4 → own group

Final Clusters:

- **Group A:** {1, 2} via email → plus 1 and c → full group = {a, b, c}
- **Group B:** {3}
- **Group C:** {4}

Select latest per group:

- Group A → Lead c (latest)
- Group B → Lead 3
- Group C → Lead 4

Summary Table:

Feature	Approach 1 (Simple Map)	Approach 2 (Union-Find)
Handles transitive duplicates	No	Yes
Time complexity	$O(n)$	$O(n * \alpha(n))$
Space complexity	$O(n)$	$O(n)$
Final entries in sample input	5	4
Accuracy of deduplication	Partial (pairwise only)	Full (transitive merge)